

Exercise 2

Tobias Raidl, 11717659

2023-10-22

Contents

Task 1	2
Task 2	5
Task 3	7
Task 4	11
Task 5	13
Task 7	14
Task 8	19

```
library("pgmm")
```

```
## Warning: package 'pgmm' was built under R version 4.1.3
```

```
library("dplyr")
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
library("ggplot2")
```

```
## Warning: package 'ggplot2' was built under R version 4.1.3
```

```
library("e1071")
```

```
## Warning: package 'e1071' was built under R version 4.1.3
```

```
library("mclust")
```

```
## Warning: package 'mclust' was built under R version 4.1.3
```

```
## Package 'mclust' version 6.0.0
```

```
## Type 'citation("mclust")' for citing this R package in publications.
```

```

library("factoextra")

## Warning: package 'factoextra' was built under R version 4.1.3
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
library("robCompositions")

## Warning: package 'robCompositions' was built under R version 4.1.3
## Loading required package: pls
##
## Attaching package: 'pls'
## The following object is masked from 'package:stats':
##
##     loadings
## Loading required package: data.table
## Warning: package 'data.table' was built under R version 4.1.3
##
## Attaching package: 'data.table'
## The following objects are masked from 'package:dplyr':
##
##     between, first, last
## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg     ggplot2
## Registered S3 method overwritten by 'perry':
##   method      from
##   print.cvFolds cvTools
library("cluster")

## Warning: package 'cluster' was built under R version 4.1.3
data(wine)
wine$Type = as.factor(wine$Type)

```


Task 1

Applying k-means on scaled (standardized) variables makes sure to avoid distorted results. This can be the case when the means or variances differ a lot between each variable. Non scaled variables cannot meaningfully be compared by distance. It is like stating 10km is less than 11m because $10 < 11$.

```

df1 = select(wine, c("Type", "Magnesium", "Total Phenols"))
ggplot(wine, aes(x=Magnesium, y=`Total Phenols`, colour=Type)) +
  geom_point()


```



```
k = 3

res.km1 = kmeans(select(df1, -Type), k, nstart=10)
df1["y_pred"] = as.factor(res.km1$cluster)

ggplot(df1, aes(x=Magnesium, y=`Total Phenols`, colour=Type)) +
  geom_point(aes(shape=y_pred))
```



```
table(df1$Type, df1$y_pred)
```

```
##
##      1  2  3
##      1 15 35  9
##      2  6 19 46
##      3  5 24 19
```

```
matchClasses(table(df1$Type, df1$y_pred))
```


```
## Cases in matched pairs: 58.99 %
```

```
## 1 2 3
## 2 3 2
```

```
df1 = select(wine, c("Type", "Magnesium", "Total Phenols"))
df1[,-c(1)] = scale(df1[,-c(1)])
```

```
res.km1.scaled = kmeans(select(df1, -Type), k, nstart=10)
df1["y_pred_scaled"] = as.factor(res.km1.scaled$cluster)
```

```
ggplot(df1, aes(x=Magnesium, y=`Total Phenols`, colour=Type)) +
  geom_point(aes(shape=y_pred_scaled))
```



```
table(df1$Type, df1$y_pred_scaled)
```

```
##
##      1  2  3
##  1 19 39  1
##  2  7 26 38
##  3  6  3 39
```

```
matchClasses(table(df1$Type, df1$y_pred_scaled))
```

```
## Cases in matched pairs: 65.17 %
```

```
## 1 2 3
## 2 3 3
```

Task 2


```
wine[, -c(1)] = scale(wine[, -c(1)])
df2 = wine

k = 3
df2.km2 = kmeans(select(df2, -Type), center=k)

df2["y_pred"] = as.factor(df2.km2$cluster)

ggplot(df2, aes(x=Magnesium, y=`Total Phenols`, colour=Type)) +
```

```
geom_point(aes(shape=y_pred))
```



```
table(df2$Type, df2$y_pred)
```

```
##
##      1  2  3
##  1 58  1  0
##  2  2 68  1
##  3  0  0 48
```

```
matchClasses(table(df2$Type, df2$y_pred))
```

```
## Cases in matched pairs: 97.75 %
```

```
## 1 2 3
## 1 2 3
```

The variables with the highest sum of pairwise distances between each types cluster center are the ones that contain most information for clustering. In our case it is Flavanoids.

```
sum_of_pairwise_distances = function(vec) {
  sum(c(abs(vec[1]-vec[2]),abs(vec[1]-vec[3]),abs(vec[2]-vec[3])))
}
```


```
df2.km.centers = data.frame(df2.km2$centers)
sum_pairwise_distance = data.frame("sum_pairwise_dist" = t(apply(df2.km.centers, 2,
  ↪ FUN=sum_of_pairwise_distances)))
```

```

most_information = colnames(sum_pairwise_distance)[apply(sum_pairwise_distance, 1,
↳ which.max)]
cat("Most information provided by variable: ", most_information)

## Most information provided by variable:  sum_pairwise_dist.Flavanoids
plot(x=wine$Type, y=wine$Flavanoids)

```



Task 3


Hierarchical clustering with methods single and centroid result in the same pair matching rate. Method complete is worse than these two methods.

```

df3 = wine
df3.dist = dist(select(df3, -Type))
df3.hc.complete = hclust(df3.dist, method="complete")
df3.hc.single = hclust(df3.dist, method="single")
df3.hc.centroid = hclust(df3.dist, method="centroid")
plot(df3.hc.complete)

```


Cluster Dendrogram



```
df3.dist  
hclust (*, "complete")
```

```
plot(df3.hc.single)
```



Cluster Dendrogram



df3.dist
hclust (*, "single")

```
plot(df3.hc.centroid)
```

Cluster Dendrogram



df3.dist
hclust (*, "centroid")

```
df3["y_pred_complete"] = as.factor(cutree(df3.hc.complete, k=3))
df3["y_pred_single"] = as.factor(cutree(df3.hc.single, k=3))
df3["y_pred_centroid"] = as.factor(cutree(df3.hc.centroid, k=3))
```

```
cat("\nComplete:\n")
```

```
##
```

```
## Complete:
```

```
matchClasses(table(df3$Type, df3$y_pred_complete))
```

```
## Cases in matched pairs: 89.89 %
```

```
## 1 2 3
```

```
## 1 2 3
```

```
cat("\nSingle:\n")
```

```
##
```

```
## Single:
```

```
matchClasses(table(df3$Type, df3$y_pred_single))
```

```
## Cases in matched pairs: 98.88 %
```

```
## 1 2 3
```

```
## 1 1 1
```

```
cat("\nCentroid:\n")

##
## Centroid:
matchClasses(table(df3$Type, df3$y_pred_centroid))

## Cases in matched pairs: 98.88 %
## 1 2 3
## 1 1 1
```

Task 4

Optimal model is VVE with num. of clusters 3.

```
df4 = wine
df4.mc = Mclust(select(df4, -Type), 1:10)
summary(df4.mc)


## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
##
## Mclust VVE (ellipsoidal, equal orientation) model with 3 components:
##
##   log-likelihood    n df         BIC          ICL
##      -4640.44 178 515  -11949.5  -11949.52
##
## Clustering table:
##   1  2  3
## 59 73 46

fviz_mclust_bic(df4.mc)


## Warning: `gather_()` was deprecated in tidyr 1.2.0.
## i Please use `gather()` instead.
## i The deprecated feature was likely used in the factoextra package.
##   Please report the issue at <https://github.com/kassambara/factoextra/issues>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

Model selection

Best model: VVE | Optimal clusters: n = 3



```
fviz_mclust(df4.mc, "classification", geom = "point")
```



```
df4["y_pred"] = as.factor(df4.mc$classification)
```

```
table(df4$Type, df4$y_pred)
```

```
##
##      1  2  3
##  1 58  1  0
##  2  1 70  0
##  3  0  2 46
```

```
matchClasses(table(df4$Type, df4$y_pred))
```

```
## Cases in matched pairs: 97.75 %
```

```
## 1 2 3
## 1 2 3
```

Task 5

The ternary dendrogram provides an overview of the cluster memberships for each observation. The color corresponds to the observations actual type. Observations of type 1 and 3 are predicted with a higher confidence than those of type 2, because their coefficients have a higher variance in the dendrogram.

```
df5 = wine
df5.cm = cmeans(df5, centers=3)
df5$y_pred = as.factor(df5.cm$cluster)
```

```
table(df5$Type, df5$y_pred)
```


```
##  
##      1  2  3  
##  1 59  0  0  
##  2  8 11 52  
##  3  0 48  0
```

```
matchClasses(table(df5$Type, df5$y_pred))
```

```
## Cases in matched pairs: 89.33 %
```

```
## 1 2 3  
## 1 3 2
```

```
ternaryDiag(df5.cm$membership, col=df5$Type)
```



Task 7

2 or 3 clusters have the lowest average silhouette width and therefore fit the requirements for optimal num. of cluster.

```
df7 = wine  
df7.dist = dist(select(df7, -Type))  
  
for(k in 2:6) {
```

```
km = unlist(kmeans(select(df7, -Type), center=k, nstart=10)$cluster)
plot(silhouette(km, df7.dist))
}
```

Silhouette plot of (x = km, dist = df7.dist)


n = 178

2 clusters C_j

$j : n_j \mid \text{ave}_{i \in C_j} s_i$

1 : 52 | 0.25


2 : 126 | 0.19



Average silhouette width : 0.21

Silhouette plot of (x = km, dist = df7.dist)


n = 178



Average silhouette width : 0.2


Silhouette plot of (x = km, dist = df7.dist)

n = 178




Silhouette plot of (x = km, dist = df7.dist)

n = 178



Silhouette plot of (x = km, dist = df7.dist)

n = 178




Task 8

The optimal num. of clusters is dependent on the correct classification - model complexity tradeoff. Lower k corresponds to lower complexity. Therefore 3 clusters seem to be the optimal fit.

```
df8 = wine
df8.cg = clusGap(select(df8, -Type), FUN=kmeans, K.max=10)
plot(df8.cg)
```

clusGap(x = select(df8, -Type), FUNcluster = kmeans, K.max = 10)



Task 9

```
df9 = wine
df9.cg <- clusGap(select(df9, -Type), FUNcluster = cmeans, K.max = 10)
plot(df9.cg)
```

**clusGap(x = select(df9, -Type), FUNcluster = cmeans, K.max
= 10)**

